


Class 7: Machine Learning 1

Morgan Black (PID A14904860)

Sample data to do clustering analysis where we know what the answer should be. Use the 'rnorm()' function.


```
#Make two individual histograms, one centered around 3 and one centered  
#around -3.  
hist( rnorm(150000, mean=3) )
```

Histogram of rnorm(150000, mean = 3)




```
hist( rnorm(150000, mean=-3) )
```

Histogram of rnorm(150000, mean = -3)




```
#Combine these two distributions on one histogram plot.  
n=15000  
hist( c(rnorm(n, mean=3), rnorm(n, mean=-3)))
```

Histogram of c(rnorm(n, mean = 3), rnorm(n, mean = -3))



```
#Make a clustering plot with rnorm() data clustered around +3 and -3
n=30
x <- c(rnorm(n, mean=3), rnorm(n, mean=-3))
y <- rev(x) #This just reverses the x axis data
z <- cbind(x,y) #Combine x and y to data matrix by column

plot(z)
```



K-means clustering

Main function for k-means clustering is 'kmeans()'


```
#Need two inputs for kmeans(): x (numeric data matrix), and centers (number of
#clusters or a set of initial cluster centers which are randomly chosen from x
#as the initial centers).

#K-means clustering with two clusters from the rnorm data above.
km <- kmeans(z, centers = 2)


#Print out the cluster membership vector (which cluster each point belongs to).
km$cluster
```

```
#Color the clusters differently by assigning the values in km$cluster vector
#to the color argument, since col=1 is black and col=2 is red.
plot(z, col=km$cluster)

#Now visualize cluster centers on the current plot:
points(km$centers, col="blue", pch=15, cex=2)
```



```
km4 <- kmeans(z, centers = 4)
plot(z, col=km4$cluster)
points(km4$centers, col="blue", pch=15, cex=2)
```



```
#This data does not look like it has 4 separate clusters, but kmeans() will
#always cluster based on what input/argument you give it.
```

Hierarchical clustering

The main function of hierarchical clustering in base R is called ‘`hclust()`’. Unlike “`kmeans()`”, we can’t just put in our data as input. First you need a distance matrix from the data


```
d <- dist(z)
hc <- hclust(d)
hc
```

```
Call:
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

There is a specific hclust plot method

```
#Plot a cluster dendrogram
plot(hc)
abline(h=10, col="red")
```




To get the main clustering result (membership vector, which cluster each point belongs to), you can cut the tree at a given height. To do this, use `'cutree()'`

```
grps <- cutree(hc, h=10)
grps
```

[illegible]

```
plot(z, col=grps)
```




Principal Component Analysis

Visualize major similarities and differences by reducing dimensionality


###PCA of UK food data

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
#head(x)


#Bar plots, not very useful
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```




```
#Pairwise plot between each pairs of countries, slightly more useful in the
#sense that all of the N. Ireland plots have more outliers compared to the
#other pairwise plots, but still hard to read.
pairs(x, col=rainbow(10), pch=16)
```



The main function to do PCA in base R is called ‘prcomp()’

```
#Need to transpose dataset to switch the rows and columns:
#t(x)

pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
#Proportion of variance tells us what portion of the original variance in the
#data is explained/captured by each PC
```

```
#What's inside the PCA object we just calculated?  
attributes(pca)
```


```
$names  
[1] "sdev"      "rotation" "center"    "scale"     "x"  
  
$class  
[1] "prcomp"
```

Make the main result figure, called a “PC plot” (or score plot, or ordination plot, or PC1 vs PC2 plot),

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
plot(pca$x[,1], pca$x[,2],  
     col=c("black", "red", "blue", "darkgreen"),  
     xlab="PC1 (67.4%)", ylab="PC2 (29%)",  
     xlim=c(-270, 500))  
text(pca$x[,1], pca$x[,2], colnames(x), col=c("black", "red", "blue", "darkgreen"))
```



```
#Variable loadings plot to see how the original variables contribute to our  
#PC1 axis since PC1 accounts for most of the variance.  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```

