

# Class 8: PCA Mini Project

Morgan Black (PID A14904860)

**Side note before starting about scaling:**

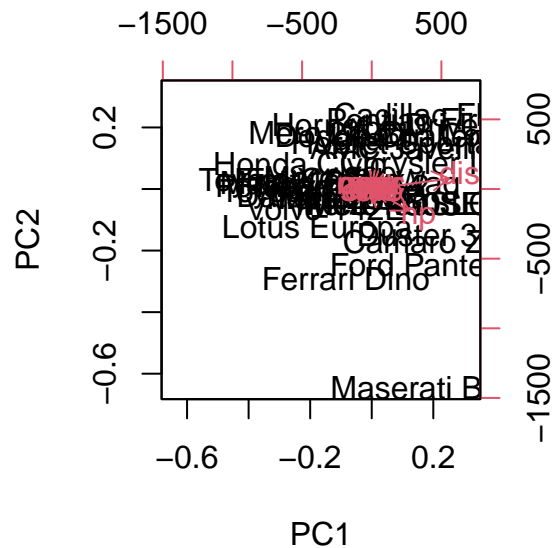
```
#Mean value of each column in mtcars dataset  
apply(mtcars, 2, mean)
```

mpg	cyl	disp	hp	drat	wt	qsec
20.090625	6.187500	230.721875	146.687500	3.596563	3.217250	17.848750
vs	am	gear	carb			
0.437500	0.406250	3.687500	2.812500			

```
#Spread of each column via standard deviation  
apply(mtcars, 2, sd)
```

mpg	cyl	disp	hp	drat	wt
6.0269481	1.7859216	123.9386938	68.5628685	0.5346787	0.9784574
qsec	vs	am	gear	carb	
1.7869432	0.5040161	0.4989909	0.7378041	1.6152000	

```
pca <- prcomp(mtcars)  
biplot(pca)
```



Without scaling, the columns are measured in different units and the pca analysis will be biased towards the units with higher counts/larger spread.

```
#Scale the data
mtscale <- scale(mtcars)
head(mtscale)
```

	mpg	cyl	disp	hp	drat
Mazda RX4	0.1508848	-0.1049878	-0.57061982	-0.5350928	0.5675137
Mazda RX4 Wag	0.1508848	-0.1049878	-0.57061982	-0.5350928	0.5675137
Datsun 710	0.4495434	-1.2248578	-0.99018209	-0.7830405	0.4739996
Hornet 4 Drive	0.2172534	-0.1049878	0.22009369	-0.5350928	-0.9661175
Hornet Sportabout	-0.2307345	1.0148821	1.04308123	0.4129422	-0.8351978
Valiant	-0.3302874	-0.1049878	-0.04616698	-0.6080186	-1.5646078

	wt	qsec	vs	am	gear
Mazda RX4	-0.610399567	-0.7771651	-0.8680278	1.1899014	0.4235542
Mazda RX4 Wag	-0.349785269	-0.4637808	-0.8680278	1.1899014	0.4235542
Datsun 710	-0.917004624	0.4260068	1.1160357	1.1899014	0.4235542
Hornet 4 Drive	-0.002299538	0.8904872	1.1160357	-0.8141431	-0.9318192
Hornet Sportabout	0.227654255	-0.4637808	-0.8680278	-0.8141431	-0.9318192
Valiant	0.248094592	1.3269868	1.1160357	-0.8141431	-0.9318192

	carb
Mazda RX4	0.7352031

```

Mazda RX4 Wag      0.7352031
Datsun 710         -1.1221521
Hornet 4 Drive     -1.1221521
Hornet Sportabout -0.5030337
Valiant            -1.1221521

```

```

#Now look at the mean and standard deviation of each column
round(apply(mtscale, 2, mean), 3)

```

```

mpg  cyl  disp    hp  drat    wt  qsec    vs    am  gear  carb
0    0    0      0    0      0    0      0    0    0     0

```

```

round(apply(mtscale, 2, sd), 3)

```

```

mpg  cyl  disp    hp  drat    wt  qsec    vs    am  gear  carb
1    1    1      1    1      1    1      1    1    1     1


```

Let's plot to make sure the scaled data still has the same relationships. Plot 'mpg' vs 'disp' for both the original and scaled data


```

library(ggplot2)
ggplot(mtcars, aes(mpg, disp)) +
  geom_point()


```



```
ggplot(mtscale, aes(mpg, disp)) +  
  geom_point()
```



```
pca2 <- prcomp(mtscale)  
biplot(pca2)
```



## Breast Cancer FNA Data

### Preparing the data

First download the csv file from the class website and put it in the current working directory thru Finder, then read it into R

```
wisc.df <- read.csv("WisconsinCancer.csv",
                    row.names=1)
#head(wisc.df)
```

Omit the 'diagnosis' column and create a new dataset

```
wisc.data <- wisc.df[, -1]
```

Create a new vector that contains the diagnosis column from the original dataset.

```
diagnosis <- as.factor(wisc.df$diagnosis)
```

**Q1: How many observations are in this dataset?**

```
nrow(wisc.df)
```

```
[1] 569
```

**Q2: How many of the observations have a malignant diagnosis?**

```
table(wisc.df$diagnosis)
```

```
  B    M  
357 212
```

There are 212 malignant diagnoses.

**Q3: How many variables/features in the data are suffixed with '\_mean'?**

```
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

There are 10.

## Principal Component Analysis

Check the mean and standard deviation of the features to determine if the data should be scaled.

```
round(colMeans(wisc.data), 3)
```

radius_mean	texture_mean	perimeter_mean
14.127	19.290	91.969
area_mean	smoothness_mean	compactness_mean
654.889	0.096	0.104
concavity_mean	concave.points_mean	symmetry_mean
0.089	0.049	0.181
fractal_dimension_mean	radius_se	texture_se
0.063	0.405	1.217
perimeter_se	area_se	smoothness_se
2.866	40.337	0.007
compactness_se	concavity_se	concave.points_se
0.025	0.032	0.012
symmetry_se	fractal_dimension_se	radius_worst
0.021	0.004	16.269
texture_worst	perimeter_worst	area_worst
25.677	107.261	880.583
smoothness_worst	compactness_worst	concavity_worst
0.132	0.254	0.272
concave.points_worst	symmetry_worst	fractal_dimension_worst
0.115	0.290	0.084

```
round(apply(wisc.data, 2, sd), 3)
```

radius_mean	texture_mean	perimeter_mean
3.524	4.301	24.299
area_mean	smoothness_mean	compactness_mean
351.914	0.014	0.053
concavity_mean	concave.points_mean	symmetry_mean
0.080	0.039	0.027
fractal_dimension_mean	radius_se	texture_se
0.007	0.277	0.552
perimeter_se	area_se	smoothness_se
2.022	45.491	0.003
compactness_se	concavity_se	concave.points_se
0.018	0.030	0.006
symmetry_se	fractal_dimension_se	radius_worst
0.008	0.003	4.833
texture_worst	perimeter_worst	area_worst
6.146	33.603	569.357
smoothness_worst	compactness_worst	concavity_worst
0.023	0.157	0.209
concave.points_worst	symmetry_worst	fractal_dimension_worst

0.066

0.062

0.018


Perform PCA on wisc.data with scaling.

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
x <- summary(wisc.pr)
x$importance
```

	PC1	PC2	PC3	PC4	PC5	PC6
Standard deviation	3.644394	2.385656	1.678675	1.407352	1.284029	1.098798
Proportion of Variance	0.442720	0.189710	0.093930	0.066020	0.054960	0.040250
Cumulative Proportion	0.442720	0.632430	0.726360	0.792390	0.847340	0.887590
	PC7	PC8	PC9	PC10	PC11	
Standard deviation	0.8217178	0.6903746	0.6456739	0.5921938	0.5421399	
Proportion of Variance	0.0225100	0.0158900	0.0139000	0.0116900	0.0098000	
Cumulative Proportion	0.9101000	0.9259800	0.9398800	0.9515700	0.9613700	
	PC12	PC13	PC14	PC15	PC16	
Standard deviation	0.5110395	0.4912815	0.3962445	0.3068142	0.2826001	
Proportion of Variance	0.0087100	0.0080500	0.0052300	0.0031400	0.0026600	
Cumulative Proportion	0.9700700	0.9781200	0.9833500	0.9864900	0.9891500	
	PC17	PC18	PC19	PC20	PC21	
Standard deviation	0.2437192	0.2293878	0.2224356	0.1765203	0.1731268	
Proportion of Variance	0.0019800	0.0017500	0.0016500	0.0010400	0.0010000	
Cumulative Proportion	0.9911300	0.9928800	0.9945300	0.9955700	0.9965700	
	PC22	PC23	PC24	PC25	PC26	
Standard deviation	0.1656484	0.1560155	0.1343689	0.1244238	0.0904303	
Proportion of Variance	0.0009100	0.0008100	0.0006000	0.0005200	0.0002700	
Cumulative Proportion	0.9974900	0.9983000	0.9989000	0.9994200	0.9996900	
	PC27	PC28	PC29	PC30		
Standard deviation	0.08306903	0.0398665	0.02736427	0.01153451		
Proportion of Variance	0.00023000	0.0000500	0.00002000	0.00000000		
Cumulative Proportion	0.99992000	0.9999700	1.00000000	1.00000000		

```
#Plot the proportion of variance accounted for by each PC
plot(x$importance[2,], typ='b')
```





**Q4: What proportion of the original variance is captured by the first PC?**

About 44% of the variance is captured by PC1.

**Q5: How many PCs are required to describe at least 70% of the original variance?**

Three principal components are required to describe at least 70% of the variance (PC3 cumulative proportion is ~72.64%).

**Q6: How many PCs are required to describe at least 90% of the original variance?**

Seven principal components are required to describe at least 90% of the variance (PC7 cumulative proportion is ~91%).


```
#biplot(wisc.pr)
```

**Q7: What stands out to you about this biplot?**

This plot is not understandable at all and is much too crowded to make any conclusions.

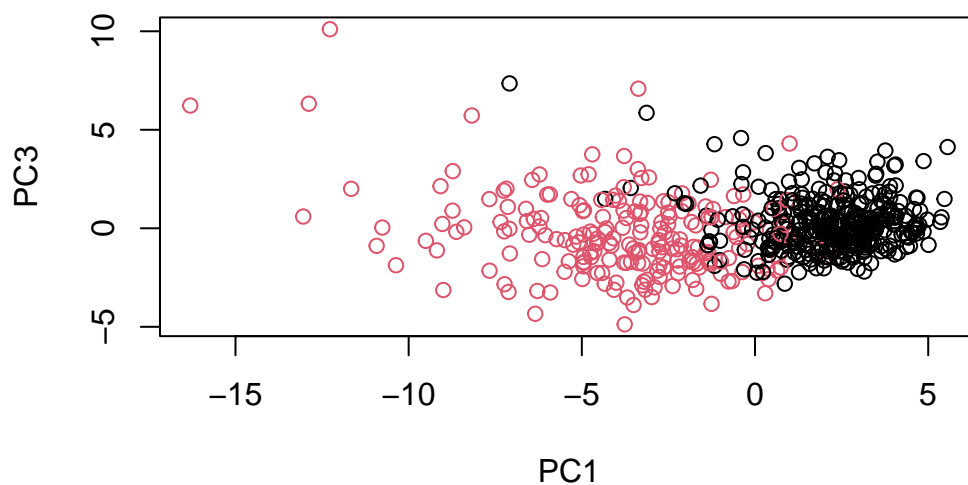
Plot the observations by PC1 and PC2:

```
#attributes(wisc.pr)
#head(wisc.pr)
plot(wisc.pr$x, col=diagnosis,
      xlab="PC1", ylab="PC2")
```



Plot the observations by PC1 and PC3:

```
plot(wisc.pr$x[,1], wisc.pr$x[,3], col=diagnosis,
      xlab="PC1", ylab="PC3")
```



**Q8: What do you notice about the two previous plots?**

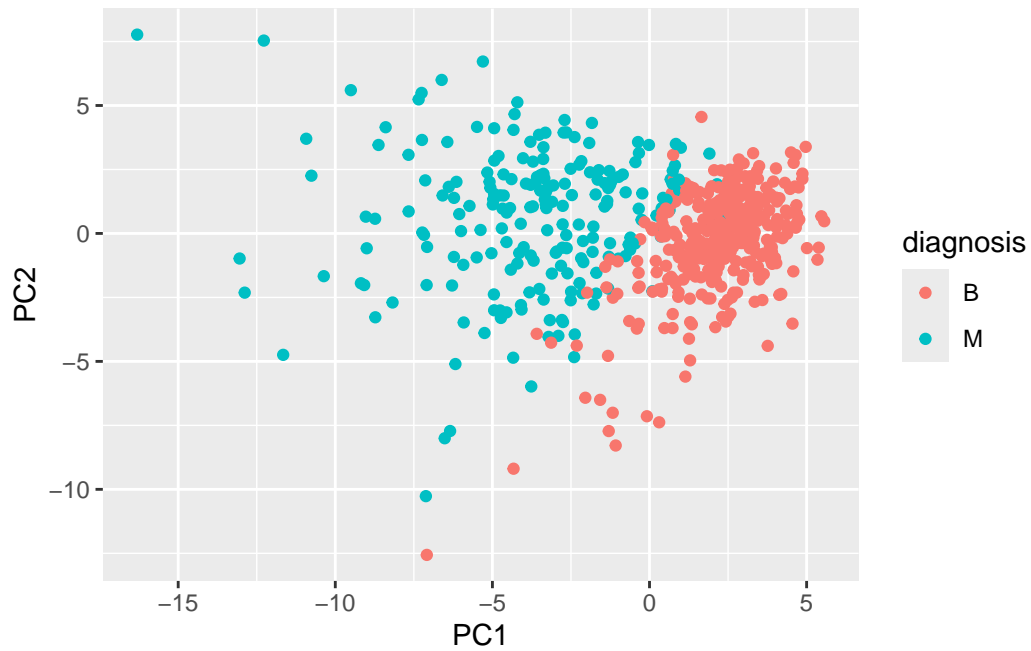
There is a cleaner more clear separation between clusters in the PC1 vs PC2 plot than the PC1 vs PC3 plot, indicating that PC2 explains more variance than PC3.

Use ggplot2 to make some nicer figures.

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

library(ggplot2)

ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```




## Variance explained

```
#Variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
#Variance explained by each PC
pve <- pr.var/sum(pr.var)

plot(pve, xlab="Principal Component",
     ylab="Proportion of Variance Explained",
     ylim=c(0,1), type="o")
```



```
#Scree plot of the same data but in a bar plot with labels for each PC
barplot(pve, ylab="Percent of Variance Explained",
        names.arg=paste0("PC", 1:length(pve)), las=2, axes=FALSE)
axis(2, at=pve, labels=round(pve,2)*100) #Creates y-axis bar label showing
```




```
#percentages rather than proportion
```

```
#ggplot based graph instead of base R bar plot
#install.packages("factoextra")
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```



## Communicating PCA results

**Q9: For the first PC, what is the component of the loading vector for the feature 'concave.points\_mean'? This tells us how much this original feature contributes to the first PC.**

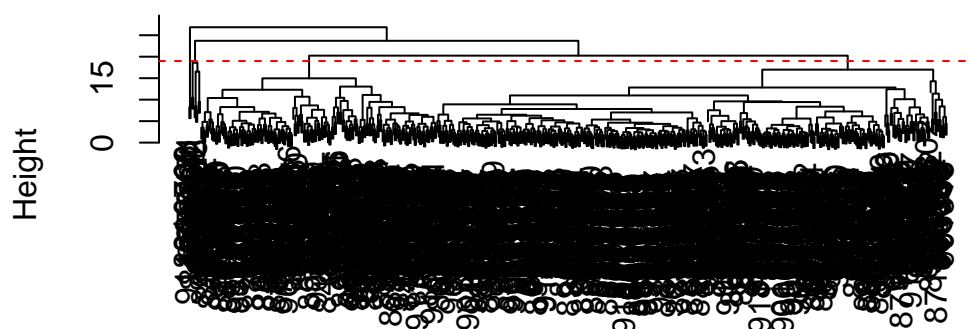
```
wisc.pr$rotation["concave.points_mean",1]
```

```
[1] -0.2608538
```

## Hierarchical clustering

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist)
plot(wisc.hclust)
abline(a=19, b=0, col="red", lty=2)
```

## Cluster Dendrogram



```
data.dist
hclust (*, "complete")
```

**Q10: What's the height where this model has 4 clusters?**

19

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis	
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Explore different methods to combine points in hierarchical clustering.

```
wisc.hclust.single <- hclust(data.dist, method="single")
wisc.hclust.cluster.single <- cutree(wisc.hclust.single, k=4)
table(wisc.hclust.cluster.single, diagnosis)
```



	diagnosis	
wisc.hclust.cluster.single	B	M
1	356	209
2	1	0
3	0	2
4	0	1

```
wisc.hclust.complete <- hclust(data.dist, method="complete")
wisc.hclust.cluster.complete <- cutree(wisc.hclust.complete, k=4)
table(wisc.hclust.cluster.complete, diagnosis)
```

	diagnosis	
wisc.hclust.cluster.complete	B	M
1	12	165
2	2	5
3	343	40
4	0	2

```
wisc.hclust.average <- hclust(data.dist, method="average")
wisc.hclust.cluster.average <- cutree(wisc.hclust.average, k=4)
table(wisc.hclust.cluster.average, diagnosis)
```

	diagnosis	
wisc.hclust.cluster.average	B	M
1	355	209
2	2	0
3	0	1
4	0	2


```
wisc.hclust.ward.D2 <- hclust(data.dist, method="ward.D2")
wisc.hclust.cluster.ward.D2 <- cutree(wisc.hclust.ward.D2, k=4)
table(wisc.hclust.cluster.ward.D2, diagnosis)
```

	diagnosis	
wisc.hclust.cluster.ward.D2	B	M
1	0	115
2	6	48
3	337	48
4	14	1

**Q12: The “complete” method seems to give the best results compared to the expert diagnoses, giving less false positive/negative results but it is unclear whether this is a good thing or just fitting the data to match the experts conclusions.**

Now use PCA results to cluster, using the “ward.D2” method

```
d<- dist(wisc.pr$x[,1:7])
hc <- hclust(d, method="ward.D2")
plot(hc)
```



Cut the tree to yield 2 clusters

```
grps <- cutree(hc, k=2)
table(grps)
```

```
grps
 1    2
216 353
```

Compare to the expert diagnoses of “M” vs “B”

```
table(diagnosis, grps)
```

	grps	
diagnosis	1	2
B	28	329
M	188	24

**Q13: How well does this model separate out the two diagnoses?**

This separates out the two diagnoses fairly well, but it is saying there are 28 false benign and 24 false malignant diagnoses from the analysis done here.

```
table(diagnosis, wisc.hclust.clusters)
```


	wisc.hclust.clusters				
diagnosis	1	2	3	4	
B	12	2	343	0	
M	165	5	40	2	

**Q14: How well does the earlier hierarchical clustering model (before PCA) do in terms of separating out the diagnoses?**


This model results in less false benign results compared to the expert diagnoses results. However, there are much more false malignant results in this model than in the PCA-based clustering model.

```
#plot(wisc.pr$x[,1:2], col=grps)

g <- as.factor(grps)
#levels(g)
g <- relevel(g,2)
#levels(g)
plot(wisc.pr$x[,1:2], col=g)
```




```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



## Prediction

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
#npc

plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



**Q16: Which of these new patient groups should we prioritize for follow up based on your results?**

Since the numeral 1 was assigned to M (malignant) and the numeral 2 was assigned to B (benign) in this analysis, the patients in cluster 1 should be prioritized for follow-up since they have been both clustered into a malignant cluster by bioinformatics analysis and diagnosed by an expert as having malignant tissue.

```
loadings <- wisc.pr$rotation
```

```
ggplot(loadings) +  
  aes(abs(PC1), reorder(rownames(loadings), -PC1)) +  
  geom_col()
```

